

# Multiprocessing using Python

CPEN333 - 2021 W2

University of British Columbia

© *Farshid Agharebparast*



# Introduction

---

- We have already studied some fundamental concepts related to processes.
- Here we look at how we can implement a program that can use two or more processes. This is one fundamental way in which we can implement **concurrency** or **parallelism**.
- In this set of slides, we focus on the Python's **multiprocessing** module to achieve parallel execution using processes.
  - ❖ This will allow truly parallel execution (given capable hardware), working around the GIL limitation.

# Objectives

---

- Learn the multiprocessing module of Python's standard library
- Use the API to implement programs that can create new processes to multi-task, join and terminate
- Understand CPython's GIL

# Let's start with an example

---

➤ First a simple single-process program:

```
def sayHello():  
    print("Hello from the function sayHello")  
  
if __name__ == "__main__":  
    print("This is the main process: beginning")  
    sayHello()  
    print("This is the main process: end")
```



Output:

```
This is the main process: beginning  
Hello from the function sayHello  
This is the main process: end
```

# Using a second process

- In the previous example, what if we want to create and delegate the execution of the sayHello method to another process (a **worker**).
- Why would we want to do that?
  - ❖ maybe we want to use the capabilities of the hardware (e.g. CPU cores) and run two things at the same time.
  - ❖ maybe we want to make our program more responsive (for example one process dealing with user interface and another with another task)
  - ❖ ...

- For example:  
(hardware and context dependant)



# Example: using a second process

---

```
import multiprocessing as mp

def sayHello():
    print("Hello from the child process")

if __name__ == "__main__":
    print("This is the main process, before starting a new process")
    p1 = mp.Process(target=sayHello)
    p1.start()
    p1.join()
    print("This is the main process after joining")
```

Output:

```
This is the main process, before starting a new process
Hello from the child process
This is the main process after joining
```

# A simple visualization

- The following simple graph shows the flow of execution in our very simple code for our parent process and the newly spawned process

Start of the parent process execution

New process spawned

Waiting  
Join()

Here we are waiting,  
but we could be doing  
other useful things

The second process finishes execution

Parent process done

# multiprocessing module

---

- **multiprocessing** is a package from Python's standard library.
- It allows us to fully leverage multiple processors (e.g. CPU cores) of a given machine.
  - ❖ To use it: `import multiprocessing`
- The following simple program will print the number of CPU's in a given system:

```
import multiprocessing

# cpu_count() returns the number of available CPUs
print("CPU count: ", multiprocessing.cpu_count())
```



# The multiprocessing.Process class

---

- We can use the **Process** class to create a process object.
- Then we use the start method to spawn the new process.
- In our previous example, we did not pass any parameters to the new process.
- Here is an example that we pass a string:

```
from multiprocessing import Process

def sayHello(name):
    print("Hello", name)

if __name__ == '__main__':
    p = Process(target=sayHello, args=("Victor",))
    p.start()
    p.join()
```

This comma  
is needed to  
indicate that  
this is a tuple

# multiprocessing.Process documentation

➤ <https://docs.python.org/3/library/multiprocessing.html#multiprocessing.Process>

```
class multiprocessing.Process(group=None, target=None, name=None, args=(), kwargs={}, *,  
daemon=None)
```

Process objects represent activity that is run in a separate process. The `Process` class has equivalents of all the methods of `threading.Thread`.

The constructor should always be called with keyword arguments. *group* should always be `None`; it exists solely for compatibility with `threading.Thread`. *target* is the callable object to be invoked by the `run()` method. It defaults to `None`, meaning nothing is called. *name* is the process name (see `name` for more details). *args* is the argument tuple for the target invocation. *kwargs* is a dictionary of keyword arguments for the target invocation. If provided, the keyword-only *daemon* argument sets the process `daemon` flag to `True` or `False`. If `None` (the default), this flag will be inherited from the creating process.

By default, no arguments are passed to *target*.

If a subclass overrides the constructor, it must make sure it invokes the base class constructor (`Process.__init__()`) before doing anything else to the process.

*Changed in version 3.3:* Added the *daemon* argument.

# start() and join()

---

- **start()**: Start the process's activity.
  - ❖ This must be called at most once per process object.
  
- **join([timeout])**:
  - ❖ If the optional argument timeout is None (the default), the method blocks until the process whose join() method is called terminates.
  - ❖ If timeout is a positive number, it blocks at most timeout seconds.

# Method of starting processes

---

- There are three distinct methods of starting processes: spawn, fork and forkserver
- **Fork**: the mechanism available and default on Unix only.
  - ❖ The fork()+exec() is historically the execution model in Unix.
  - ❖ Since python version 3.8, macOS does not use fork (considered unsafe).
- **Forkserver**: This method is only available for select Unix platform.
- So we need to mainly focus on the **spawn** method.

# Spawn

---

- **Spawning** creates a second distinct python interpreter process, including its own GIL (we will discuss GIL soon).
  - ❖ Starting the process activity by calling `start()`
  - ❖ Spawn is default on Windows and macOS (since version 3.8)
- Why is this important?
  - ❖ A new process created using spawn will begin execution from the top of the program.
  - ❖ So we **need to use** a `if __name__=="__main__":` block to specify the code intended for the main process.
  - ❖ Failing to do so causes *RuntimeError* or may cause an infinitum creation of process recursively.

# Daemon Processes

---

- A **daemon** process continues to run until done or until the main process is terminated (whichever occurs first).
  - ❖ A note from the documentation: "... these are **not** Unix daemons or services, they are normal processes that will be terminated (and not joined) if non-daemonic processes have exited."
- We can use the daemon flag to daemonize a process (flag set to True)

```
dp = multiprocessing.Process(target = daemonProcess)  
dp.daemon = True
```

  - ❖ The flag must be set before `start()` is called. Alternatively, we can use the `daemon` argument of the `Process` constructor.

# Example d1

```
import multiprocessing as mp
import time

def myProcess():
    print("Starting the child Process")
    print(f"Child process started: {mp.current_process()} pid = {mp.current_process().pid}")
    time.sleep(3)
    print("Child process terminating ...")

if __name__ == "__main__":
    print(f"Main process: {mp.current_process()} pid = {mp.current_process().pid}")
    p = mp.Process(target=myProcess, name="child")
    p.daemon = True
    p.start()
    print("Let's see if the child process continues to execute")
    time.sleep(6)
    print("Main process terminating ...")
```

## Notes:

- We are just waiting using `sleep` in main (e.g. instead of `join`) in order to be able to examine these different scenarios.
- Instead of `multiprocessing.current_process().pid`, we could use `os.getpid()` (from the `os` module)

# Example d1 output

---

➤ Note that:

- ❖ The child process is a daemon process
- ❖ Due to sleep, the main process waits long enough for the child process to complete and terminate, before it itself terminating.

➤ A sample output from the program is:

```
Main process: <_MainProcess name='MainProcess' parent=None started> pid = 7780
Let's see if the child process continues to execute
Starting the child Process
Child process started: <Process name='child' parent=7780 started daemon> pid = 4176
Child process terminating ...
Main process terminating ...
```



# Example d2

---

- Example d2 is similar to Example d1, except for the duration of sleep in the main process (marked below):

```
import multiprocessing as mp
import time

def myProcess():
    print("Starting the child Process")
    print(f"Child process started: {mp.current_process()} pid = {mp.current_process().pid}")
    time.sleep(3)
    print("Child process terminating ...")

if __name__ == "__main__":
    print(f"Main process: {mp.current_process()} pid = {mp.current_process().pid}")
    p = mp.Process(target=myProcess, name="child")
    p.daemon = True
    p.start()
    print("Let's see if the child process continues to execute")
    time.sleep(1)
    print("Main process terminating ...")
```

# Example d2 output

---

➤ Note that:

- ❖ The child process is a daemon process
- ❖ The main process does not wait long enough for the child process to complete and terminate.
- ❖ So the main process terminates before the child complete and the child process is terminated prematurely.

➤ A sample output from the program is:

```
Main process: <_MainProcess name='MainProcess' parent=None started> pid = 9056
```

```
Let's see if the child process continues to execute
```

```
Starting the child Process
```

```
Child process started: <Process name='child' parent=9056 started daemon> pid = 2444
```

```
Main process terminating ...
```

# Example d3

➤ Now let's compare example d2 with a non-daemonic child process:

```
import multiprocessing as mp
import time

def myProcess():
    print("Starting the child Process")
    print(f"Child process started: {mp.current_process()} pid = {mp.current_process().pid}")
    time.sleep(3)
    print(f"Child process (parent pid = {mp.parent_process().pid}) terminating ...")

if __name__ == "__main__":
    print(f"Main process: {mp.current_process()} pid = {mp.current_process().pid}")
    p = mp.Process(target=myProcess, name="child")
    p.daemon = False
    p.start()
    print("Let's see if the child process continues to execute")
    time.sleep(1)
    print("Main process terminating ...")
```

# Example d3 output

---

## ➤ Note that:

- ❖ The child process is not a daemon process (non-daemonic)
- ❖ So even though the main process is done before the child process completes, the child process continues and completes before terminating.
- ❖ This is the default for non-daemonic child processes, where the main program will not exit until all the children have exited (multiprocessing tries to make sure that programs using it behave well).

## ➤ A sample output from the program is:

```
Main process: <_MainProcess name='MainProcess' parent=None started> pid = 15020
Let's see if the child process continues to execute
Starting the child Process
Child process started: <Process name='child' parent=15020 started> pid = 7292
Main process terminating ...
Child process (parent pid = 15020) terminating ...
```

# Terminating a process

---

- We can use the *terminate()* method on our process object to terminate it.
- <https://docs.python.org/3/library/multiprocessing.html#multiprocessing.Process.terminate>

## **terminate()**

Terminate the process. On Unix this is done using the `SIGTERM` signal; on Windows `TerminateProcess()` is used. Note that exit handlers and finally clauses, etc., will not be executed.

Note that descendant processes of the process will *not* be terminated – they will simply become orphaned.

# Example: terminate()

```
import multiprocessing
import time

def myProcess():
    print(f"Starting the child Process with pid: {multiprocessing.current_process().pid}")
    time.sleep(10)
    print("Child process terminating normally")

if __name__ == "__main__":
    p = multiprocessing.Process(target=myProcess, name="child")
    p.start()
    time.sleep(1)
    print("Main is to terminate the child process")
    p.terminate()
    print("Child process killed by main")
```

Output:

```
Starting the child Process with pid: 23344
Main is to terminate the child process
Child process killed by main
```

# Process Pool

- The **Pool** class can be used to manage a fixed number of child processes for simple cases where the work to be done can be broken up and distributed independently.

- ❖ <https://docs.python.org/3/library/multiprocessing.html#multiprocessing.pool.Pool>

## Process Pools

One can create a pool of processes which will carry out tasks submitted to it with the `Pool` class.

```
class multiprocessing.pool.Pool([processes[, initializer[, initargs[, maxtasksperchild[, context]]]])
```

A process pool object which controls a pool of worker processes to which jobs can be submitted. It supports asynchronous results with timeouts and callbacks and has a parallel map implementation.

`processes` is the number of worker processes to use. If `processes` is `None` then the number returned by `os.cpu_count()` is used.

If `initializer` is not `None` then each worker process will call `initializer(*initargs)` when it starts.

`maxtasksperchild` is the number of tasks a worker process can complete before it will exit and be replaced with a fresh worker process, to enable unused resources to be freed. The default `maxtasksperchild` is `None`, which means worker processes will live as long as the pool.

`context` can be used to specify the context used for starting the worker processes. Usually a pool is created using the function `multiprocessing.Pool()` or the `Pool()` method of a context object. In both cases `context` is set appropriately.

# Example

---

- In this example, a pool of 4 processes are created and jobs are passed to them until there are no more jobs.
- `map()` is used here on the iterable `range(10)`. It chops the iterable into a number of chunks which it submits to the process pool as separate tasks.

```
from multiprocessing import Pool

def f(x):
    return x*x

if __name__ == '__main__':
    with Pool(processes=4) as pool:
        print(pool.map(f, range(10)))
```

# start 4 worker processes  
# prints "[0, 1, 4, ..., 81]"



# The `with` statement

---

- The `with` statement is used for resource management and exception handling.

- This generic `with` statement:

```
with some_resource:  
    # do something...
```

- is equivalent to the following code:

```
some_resource.acquire()  
try:  
    # do something...  
finally:  
    some_resource.release()
```

- <https://docs.python.org/3/library/threading.html#using-locks-conditions-and-semaphores-in-the-with-statement>

# Related topics

---

- We postpone discussions on some related topics and modules to near future, for example:
  - ❖ Topics:
    - Inter-process communication
    - Synchronization
    - Threads, ...
  - ❖ Python library:
    - threading
    - os
    - subprocess, ...

# Working around GIL

---

- The very popular CPython implementation of Python uses GIL (**global interpreter lock**).
  - ❖ We will talk about GIL and its limitation when we talk about threading next in more details.
- For now, just a note that by using multiprocessing we effectively work around GIL.
  - ❖ This way we can create multiple processes that would allow us to fully leverage multiple processors power on a machine.
  - ❖ This though will cause other challenges such as dealing with inter-process communications ...

# References

---

- Python documentation
  - ❖ <https://docs.python.org/3/library/multiprocessing.html>
- *The Python Standard Library*, D. Hellmann
- *Learning Concurrency in Python*
  - ❖ Note that this book seems to have errata. Use it with a grain of salt.